

✓ Example 1-Multiple_Linear_Regression

```
import numpy as np
from sklearn.linear_model import LinearRegression
```

```
x = [[0, 1], [5, 1], [15, 2], [25, 5], [35, 11], [45, 15], [55, 34], [60, 35]]
y = [4, 5, 20, 14, 32, 22, 38, 43]
x, y = np.array(x), np.array(y)
print(x)
print(y)
x.shape
y.shape
```

```
[[ 0  1]
 [ 5  1]
 [15  2]
 [25  5]
 [35 11]
 [45 15]
 [55 34]
 [60 35]]
[ 4  5 20 14 32 22 38 43]
(8,)
```

```
model = LinearRegression().fit(x, y)
```

```
r_sq = model.score(x, y)
print(f"coefficient of determination: {r_sq}")
print(f"intercept: {model.intercept_}")
print(f"coefficients: {model.coef_}")
```

```
coefficient of determination: 0.8615939258756776
intercept: 5.52257927519819
coefficients: [0.44706965 0.25502548]
```

```
y_pred = model.predict(x)
print(f"predicted response:\n{y_pred}")
```

```
predicted response:
[ 5.77760476  8.012953 12.73867497 17.9744479 23.97529728 29.4660957
 38.78227633 41.27265006]
```

```
# Verification of predict expression
y_pred = model.intercept_ + np.sum(model.coef_ * x, axis=1)
print(f"predicted response:\n{y_pred}")
```

```
predicted response:
[ 5.77760476  8.012953 12.73867497 17.9744479 23.97529728 29.4660957
 38.78227633 41.27265006]
```

✓ Example 2-Multiple_Linear_Regression using boston data

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
boston_data_frame= pd.read_csv("/content/drive/MyDrive/AI Tools Lab/boston_house_prices.csv")
boston_data_frame.head()
```

```
CRIM  ZN  INDUS  CHAS  NOX   RM   AGE    DIS  RAD  TAX  PTRATIO    B  LST.
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900   1  296    15.3  396.90  4.
1  0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671   2  242    17.8  396.90  9.
2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671   2  242    17.8  392.83  4.
3  0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622   3  222    18.7  394.63  2.
4  0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622   3  222    18.7  396.90  5.
```

```

"""CRIM - per capita crime rate by town
ZN - proportion of residential land zoned for lots over 25,000 sq.ft.
INDUS - proportion of non-retail business acres per town.
CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)
NOX - nitric oxides concentration (parts per 10 million)
RM - average number of rooms per dwelling
AGE - proportion of owner-occupied units built prior to 1940
DIS - weighted distances to five Boston employment centres
RAD - index of accessibility to radial highways
TAX - full-value property-tax rate per $10,000
PTRATIO - pupil-teacher ratio by town
B - 1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
LSTAT - % lower status of the population
MEDV - Median value of owner-occupied homes in $1000's
Price-Price"""

```

```

#First, check for missing information.
boston_data_frame.isnull().sum()

```

```

→ CRIM      0
   ZN       0
   INDUS    0
   CHAS     0
   NOX      0
   RM       0
   AGE      0
   DIS      0
   RAD      0
   TAX      0
   PTRATIO  0
   B        0
   LSTAT    0
   MEDV     0
   PRICE    0
dtype: int64

```

```
boston_data_frame.info()
```

```

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS        506 non-null    int64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    int64
 9   TAX         506 non-null    int64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  MEDV        506 non-null    float64
14  PRICE       506 non-null    float64
dtypes: float64(12), int64(3)
memory usage: 59.4 KB

```

```
#We can separate dependent and independent variables
```

```

boston_data_X = boston_data_frame[boston_data_frame.columns[0:13]]
boston_data_Y = boston_data_frame[boston_data_frame.columns[13:14]]

```

```

print(boston_data_X)
print(boston_data_Y)

```

```

→
   CRIM  ZN  INDUS  CHAS  NOX  RM  AGE  DIS  RAD  TAX  \
0  0.00632  18.0  2.31    0  0.538  6.575  65.2  4.0900  1  296
1  0.02731  0.0  7.07    0  0.469  6.421  78.9  4.9671  2  242
2  0.02729  0.0  7.07    0  0.469  7.185  61.1  4.9671  2  242
3  0.03237  0.0  2.18    0  0.458  6.998  45.8  6.0622  3  222
4  0.06905  0.0  2.18    0  0.458  7.147  54.2  6.0622  3  222
..  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
501 0.06263  0.0  11.93    0  0.573  6.593  69.1  2.4786  1  273
502 0.04527  0.0  11.93    0  0.573  6.120  76.7  2.2875  1  273
503 0.06076  0.0  11.93    0  0.573  6.976  91.0  2.1675  1  273
504 0.10959  0.0  11.93    0  0.573  6.794  89.3  2.3889  1  273
505 0.04741  0.0  11.93    0  0.573  6.030  80.8  2.5050  1  273

   PTRATIO  B  LSTAT
0      15.3 396.90  4.98
1      17.8 396.90  9.14

```

```

2      17.8  392.83  4.03
3      18.7  394.63  2.94
4      18.7  396.90  5.33
..      ...      ...
501    21.0  391.99  9.67
502    21.0  396.90  9.08
503    21.0  396.90  5.64
504    21.0  393.45  6.48
505    21.0  396.90  7.88

```

```
[506 rows x 13 columns]
```

```

MEDV
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
..      ...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9

```

```
[506 rows x 1 columns]
```

```

#Check the correlation
boston_data_X.corr()

```



	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.37
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.66
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.70
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.09
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.76
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.20
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.74
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.00
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.49
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.53
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.23
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.29
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.49


```

#Visualize correlation between attributes by using heatmap
sns.heatmap(boston_data_X.corr())

```



<Axes: >



```
#Feature contains high correlation. We need to remove them first before applying regression techniques. Create correlation matrix
abs_corr_matrix = boston_data_X.corr().abs()
#print(abs_corr_matrix)
#Select upper triangle of matrix
up_tri=abs_corr_matrix.where(np.triu(np.ones(abs_corr_matrix.shape),k=1).astype(np.bool_))
print(up_tri)
# Find all the features which is having correlation > 0.75 with other features.
correlated_features = [column for column in up_tri.columns if any(up_tri[column] > 0.75)]
#Print correlated_features
print(correlated_features)
```

```
CRIM      CRIM      ZN      INDUS      CHAS      NOX      RM      AGE \
CRIM      NaN      0.200469  0.406583  0.055892  0.420972  0.219247  0.352734
ZN         NaN      NaN      0.533828  0.042697  0.516604  0.311991  0.569537
INDUS      NaN      NaN      NaN      0.062938  0.763651  0.391676  0.644779
CHAS      NaN      NaN      NaN      NaN      0.091203  0.091251  0.086518
NOX      NaN      NaN      NaN      NaN      NaN      0.302188  0.731470
RM         NaN      NaN      NaN      NaN      NaN      NaN      0.240265
AGE      NaN      NaN      NaN      NaN      NaN      NaN      NaN
DIS      NaN      NaN      NaN      NaN      NaN      NaN      NaN
RAD      NaN      NaN      NaN      NaN      NaN      NaN      NaN
TAX      NaN      NaN      NaN      NaN      NaN      NaN      NaN
PTRATIO   NaN      NaN      NaN      NaN      NaN      NaN      NaN
B         NaN      NaN      NaN      NaN      NaN      NaN      NaN
LSTAT     NaN      NaN      NaN      NaN      NaN      NaN      NaN

DIS      RAD      TAX      PTRATIO      B      LSTAT
CRIM      0.379670  0.625505  0.582764  0.289946  0.385064  0.455621
ZN         0.664408  0.311948  0.314563  0.391679  0.175520  0.412995
INDUS      0.708027  0.595129  0.720760  0.383248  0.356977  0.603800
CHAS      0.099176  0.007368  0.035587  0.121515  0.048788  0.053929
NOX      0.769230  0.611441  0.668023  0.188933  0.380051  0.590879
RM         0.205246  0.209847  0.292048  0.355501  0.128069  0.613808
AGE      0.747881  0.456022  0.506456  0.261515  0.273534  0.602339
DIS      NaN      0.494588  0.534432  0.232471  0.291512  0.496996
RAD      NaN      NaN      0.910228  0.464741  0.444413  0.488676
TAX      NaN      NaN      NaN      0.460853  0.441808  0.543993
PTRATIO   NaN      NaN      NaN      NaN      0.177383  0.374044
B         NaN      NaN      NaN      NaN      NaN      0.366087
LSTAT     NaN      NaN      NaN      NaN      NaN      NaN

['NOX', 'DIS', 'TAX']
```

```
#Drop correlated features to avoid redundancy
boston_data_X = boston_data_X.drop(correlated_features,
axis=1)
```

```
#Divide the data into training and test set. Train set contains 80% of the data. Testset contains 20% of the data.
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(boston_data_X, boston_data_Y, test_size=0.20)
```


```
from sklearn.linear_model import LinearRegression
linear_regression = LinearRegression()
```

```
linear_regression.fit(X_train,Y_train)
```

```
LinearRegression
LinearRegression()
```

```
Y_pred = linear_regression.predict(X_test)
```

```
plt.scatter(Y_test, Y_pred, c = 'green')
plt.xlabel("Price")
plt.ylabel("Predicted value")
plt.title("True value vs predicted value")
plt.show()
```



```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
mse = mean_squared_error(Y_test, Y_pred)
mae = mean_absolute_error(Y_test, Y_pred)
r2=r2_score(Y_test, Y_pred)
print("Mean Square Error : ", mse)
print("Mean Absolute Error : ", mae)
#MSE is more sensitive to outliers than MAE because the errors are squared, which means that larger errors have a disproportionately lar
print("R2score : ", r2)
```



```
Mean Square Error : 35.936830198059006
Mean Absolute Error : 4.005302201237243
R2score : 0.6844594894233893
```

Start coding or [generate](#) with AI.